

# Lab 10

**Aim:** To implement ResNet101 CNN Architecture

[illegible]

```

download=True)

train_dataset = Subset(train_and_valid, train_indices)
valid_dataset = Subset(train_and_valid, valid_indices)

test_dataset = datasets.CIFAR10(root='data',
                                train=False,
                                transform=transforms.ToTensor())

train_loader = DataLoader(dataset=train_dataset,
                           batch_size=BATCH_SIZE,
                           num_workers=8,
                           shuffle=True)

valid_loader = DataLoader(dataset=valid_dataset,
                           batch_size=BATCH_SIZE,
                           num_workers=8,
                           shuffle=False)

test_loader = DataLoader(dataset=test_dataset,
                           batch_size=BATCH_SIZE,
                           num_workers=8,
                           shuffle=False)

```

*# Checking the dataset*

```

for images, labels in train_loader:
    print('Image batch dimensions:', images.shape)
    print('Image label dimensions:', labels.shape)
    break

for images, labels in test_loader:
    print('Image batch dimensions:', images.shape)
    print('Image label dimensions:', labels.shape)
    break

for images, labels in valid_loader:
    print('Image batch dimensions:', images.shape)
    print('Image label dimensions:', labels.shape)
    break

```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to data/cifar-10-python.tar.gz

100%|██████████| 170498071/170498071 [00:13<00:00, 12606907.54it/s]

Extracting data/cifar-10-python.tar.gz to data

/usr/local/lib/python3.10/dist-packages/torch/utils/data/  
data\_loader.py:558: UserWarning: This DataLoader will create 8 worker

processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```
warnings.warn(_create_warning_msg(
```

```
Image batch dimensions: torch.Size([128, 3, 32, 32])
```

```
Image label dimensions: torch.Size([128])
```

```
Image batch dimensions: torch.Size([128, 3, 32, 32])
```

```
Image label dimensions: torch.Size([128])
```

```
Image batch dimensions: torch.Size([128, 3, 32, 32])
```

```
Image label dimensions: torch.Size([128])
```

```
def conv3x3(in_planes, out_planes, stride=1):  
    """3x3 convolution with padding"""  
    return nn.Conv2d(in_planes, out_planes, kernel_size=3,  
stride=stride,  
padding=1, bias=False)
```

```
class Bottleneck(nn.Module):  
    expansion = 4
```

```
    def __init__(self, inplanes, planes, stride=1, downsample=None):  
        super(Bottleneck, self).__init__()  
        self.conv1 = nn.Conv2d(inplanes, planes, kernel_size=1,  
bias=False)  
        self.bn1 = nn.BatchNorm2d(planes)  
        self.conv2 = nn.Conv2d(planes, planes, kernel_size=3,  
stride=stride,  
padding=1, bias=False)  
        self.bn2 = nn.BatchNorm2d(planes)  
        self.conv3 = nn.Conv2d(planes, planes * 4, kernel_size=1,  
bias=False)  
        self.bn3 = nn.BatchNorm2d(planes * 4)  
        self.relu = nn.ReLU(inplace=True)  
        self.downsample = downsample  
        self.stride = stride
```

```
    def forward(self, x):  
        residual = x  
  
        out = self.conv1(x)  
        out = self.bn1(out)  
        out = self.relu(out)  
  
        out = self.conv2(out)  
        out = self.bn2(out)  
        out = self.relu(out)
```

```

        out = self.conv3(out)
        out = self.bn3(out)

        if self.downsample is not None:
            residual = self.downsample(x)

        out += residual
        out = self.relu(out)

        return out

class ResNet(nn.Module):

    def __init__(self, block, layers, num_classes, grayscale):
        self.inplanes = 64
        if grayscale:
            in_dim = 1
        else:
            in_dim = 3
        super(ResNet, self).__init__()
        self.conv1 = nn.Conv2d(in_dim, 64, kernel_size=7, stride=2,
padding=3,
                                bias=False)
        self.bn1 = nn.BatchNorm2d(64)
        self.relu = nn.ReLU(inplace=True)
        self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2,
padding=1)
        self.layer1 = self._make_layer(block, 64, layers[0])
        self.layer2 = self._make_layer(block, 128, layers[1],
stride=2)
        self.layer3 = self._make_layer(block, 256, layers[2],
stride=2)
        self.layer4 = self._make_layer(block, 512, layers[3],
stride=2)
        self.avgpool = nn.AvgPool2d(7, stride=1, padding=2)
        #self.fc = nn.Linear(2048 * block.expansion, num_classes)
        self.fc = nn.Linear(2048, num_classes)

        for m in self.modules():
            if isinstance(m, nn.Conv2d):
                n = m.kernel_size[0] * m.kernel_size[1] *
m.out_channels
                m.weight.data.normal_(0, (2. / n)**.5)
            elif isinstance(m, nn.BatchNorm2d):
                m.weight.data.fill_(1)
                m.bias.data.zero_()

    def _make_layer(self, block, planes, blocks, stride=1):

```

```

        downsample = None
        if stride != 1 or self.inplanes != planes * block.expansion:
            downsample = nn.Sequential(
                nn.Conv2d(self.inplanes, planes * block.expansion,
                           kernel_size=1, stride=stride, bias=False),
                nn.BatchNorm2d(planes * block.expansion),
            )

        layers = []
        layers.append(block(self.inplanes, planes, stride,
downsample))
        self.inplanes = planes * block.expansion
        for i in range(1, blocks):
            layers.append(block(self.inplanes, planes))

        return nn.Sequential(*layers)

def forward(self, x):
    x = self.conv1(x)
    x = self.bn1(x)
    x = self.relu(x)
    x = self.maxpool(x)

    x = self.layer1(x)
    x = self.layer2(x)
    x = self.layer3(x)
    x = self.layer4(x)

    #x = self.avgpool(x)
    x = x.view(x.size(0), -1)
    logits = self.fc(x)
    probas = F.softmax(logits, dim=1)
    return logits, probas

def resnet101(num_classes, grayscale):
    """Constructs a ResNet-101 model."""
    model = ResNet(block=Bottleneck,
                    layers=[3, 4, 23, 3],
                    num_classes=NUM_CLASSES,
                    grayscale=grayscale)

    return model

torch.manual_seed(RANDOM_SEED)
model = resnet101(NUM_CLASSES, GRAYSCALE)
model.to(DEVICE)

optimizer = torch.optim.Adam(model.parameters(), lr=LEARNING_RATE)

```

```

def compute_accuracy(model, data_loader, device):
    correct_pred, num_examples = 0, 0
    for i, (features, targets) in enumerate(data_loader):

        features = features.to(device)
        targets = targets.to(device)

        logits, probas = model(features)
        _, predicted_labels = torch.max(probas, 1)
        num_examples += targets.size(0)
        correct_pred += (predicted_labels == targets).sum()
    return correct_pred.float()/num_examples * 100

start_time = time.time()

# use random seed for reproducibility (here batch shuffling)
torch.manual_seed(RANDOM_SEED)

for epoch in range(NUM_EPOCHS):

    model.train()

    for batch_idx, (features, targets) in enumerate(train_loader):

        ### PREPARE MINIBATCH
        features = features.to(DEVICE)
        targets = targets.to(DEVICE)

        ### FORWARD AND BACK PROP
        logits, probas = model(features)
        cost = F.cross_entropy(logits, targets)
        optimizer.zero_grad()

        cost.backward()

        ### UPDATE MODEL PARAMETERS
        optimizer.step()

        ### LOGGING
        if not batch_idx % 120:
            print (f'Epoch: {epoch+1:03d}/{NUM_EPOCHS:03d} | '
                  f'Batch {batch_idx:03d}/{len(train_loader):03d} | '
                  f' Cost: {cost:.4f}')

        # no need to build the computation graph for backprop when
computing accuracy
        with torch.set_grad_enabled(False):
            train_acc = compute_accuracy(model, train_loader,
            device=DEVICE)
            valid_acc = compute_accuracy(model, valid_loader,

```

```

device=DEVICE)
    print(f'Epoch: {epoch+1:03d}/{NUM_EPOCHS:03d} Train Acc.:
{train_acc:.2f}%')
        f' | Validation Acc.: {valid_acc:.2f}%')

    elapsed = (time.time() - start_time)/60
    print(f'Time elapsed: {elapsed:.2f} min')

elapsed = (time.time() - start_time)/60
print(f'Total Training Time: {elapsed:.2f} min')

```

```

Epoch: 001/025 | Batch 000/383 | Cost: 2.7585
Epoch: 001/025 | Batch 120/383 | Cost: 3.7965
Epoch: 001/025 | Batch 240/383 | Cost: 2.2732
Epoch: 001/025 | Batch 360/383 | Cost: 2.0641
Epoch: 001/025 Train Acc.: 29.08% | Validation Acc.: 30.80%
Time elapsed: 1.30 min
Epoch: 002/025 | Batch 000/383 | Cost: 1.8997
Epoch: 002/025 | Batch 120/383 | Cost: 1.6565
Epoch: 002/025 | Batch 240/383 | Cost: 1.6737
Epoch: 002/025 | Batch 360/383 | Cost: 1.7384
Epoch: 002/025 Train Acc.: 41.41% | Validation Acc.: 41.20%
Time elapsed: 2.57 min
Epoch: 003/025 | Batch 000/383 | Cost: 1.4844
Epoch: 003/025 | Batch 120/383 | Cost: 1.5908
Epoch: 003/025 | Batch 240/383 | Cost: 1.4972
Epoch: 003/025 | Batch 360/383 | Cost: 1.4840
Epoch: 003/025 Train Acc.: 48.25% | Validation Acc.: 49.10%
Time elapsed: 3.84 min
Epoch: 004/025 | Batch 000/383 | Cost: 1.1930
Epoch: 004/025 | Batch 120/383 | Cost: 1.4712
Epoch: 004/025 | Batch 240/383 | Cost: 1.4109
Epoch: 004/025 | Batch 360/383 | Cost: 1.2905
Epoch: 004/025 Train Acc.: 55.51% | Validation Acc.: 53.20%
Time elapsed: 5.13 min
Epoch: 005/025 | Batch 000/383 | Cost: 1.2403
Epoch: 005/025 | Batch 120/383 | Cost: 1.1947
Epoch: 005/025 | Batch 240/383 | Cost: 1.3734
Epoch: 005/025 | Batch 360/383 | Cost: 1.0939
Epoch: 005/025 Train Acc.: 61.72% | Validation Acc.: 59.40%
Time elapsed: 6.41 min
Epoch: 006/025 | Batch 000/383 | Cost: 1.0849
Epoch: 006/025 | Batch 120/383 | Cost: 1.0589
Epoch: 006/025 | Batch 240/383 | Cost: 1.4794
Epoch: 006/025 | Batch 360/383 | Cost: 1.1895
Epoch: 006/025 Train Acc.: 61.20% | Validation Acc.: 59.20%
Time elapsed: 7.69 min
Epoch: 007/025 | Batch 000/383 | Cost: 1.1170
Epoch: 007/025 | Batch 120/383 | Cost: 1.0913
Epoch: 007/025 | Batch 240/383 | Cost: 1.1495

```

Epoch: 007/025 | Batch 360/383 | Cost: 0.9294  
Epoch: 007/025 Train Acc.: 68.34% | Validation Acc.: 66.30%  
Time elapsed: 8.98 min  
Epoch: 008/025 | Batch 000/383 | Cost: 0.8882  
Epoch: 008/025 | Batch 120/383 | Cost: 0.9053  
Epoch: 008/025 | Batch 240/383 | Cost: 1.0006  
Epoch: 008/025 | Batch 360/383 | Cost: 0.7615  
Epoch: 008/025 Train Acc.: 74.02% | Validation Acc.: 70.10%  
Time elapsed: 10.27 min  
Epoch: 009/025 | Batch 000/383 | Cost: 0.7237  
Epoch: 009/025 | Batch 120/383 | Cost: 1.0414  
Epoch: 009/025 | Batch 240/383 | Cost: 0.8668  
Epoch: 009/025 | Batch 360/383 | Cost: 0.7295  
Epoch: 009/025 Train Acc.: 78.16% | Validation Acc.: 71.40%  
Time elapsed: 11.56 min  
Epoch: 010/025 | Batch 000/383 | Cost: 0.7486  
Epoch: 010/025 | Batch 120/383 | Cost: 0.8835  
Epoch: 010/025 | Batch 240/383 | Cost: 0.7277  
Epoch: 010/025 | Batch 360/383 | Cost: 0.8684  
Epoch: 010/025 Train Acc.: 74.77% | Validation Acc.: 69.10%  
Time elapsed: 12.84 min  
Epoch: 011/025 | Batch 000/383 | Cost: 0.6268  
Epoch: 011/025 | Batch 120/383 | Cost: 0.6856  
Epoch: 011/025 | Batch 240/383 | Cost: 0.7181  
Epoch: 011/025 | Batch 360/383 | Cost: 0.5709  
Epoch: 011/025 Train Acc.: 78.31% | Validation Acc.: 72.00%  
Time elapsed: 14.12 min  
Epoch: 012/025 | Batch 000/383 | Cost: 0.6416  
Epoch: 012/025 | Batch 120/383 | Cost: 0.7250  
Epoch: 012/025 | Batch 240/383 | Cost: 0.7191  
Epoch: 012/025 | Batch 360/383 | Cost: 0.6804  
Epoch: 012/025 Train Acc.: 81.89% | Validation Acc.: 72.30%  
Time elapsed: 15.41 min  
Epoch: 013/025 | Batch 000/383 | Cost: 0.5234  
Epoch: 013/025 | Batch 120/383 | Cost: 0.7396  
Epoch: 013/025 | Batch 240/383 | Cost: 0.5967  
Epoch: 013/025 | Batch 360/383 | Cost: 0.8753  
Epoch: 013/025 Train Acc.: 81.08% | Validation Acc.: 74.80%  
Time elapsed: 16.70 min  
Epoch: 014/025 | Batch 000/383 | Cost: 0.4325  
Epoch: 014/025 | Batch 120/383 | Cost: 0.5298  
Epoch: 014/025 | Batch 240/383 | Cost: 0.5962  
Epoch: 014/025 | Batch 360/383 | Cost: 0.6495  
Epoch: 014/025 Train Acc.: 85.73% | Validation Acc.: 75.50%  
Time elapsed: 17.99 min  
Epoch: 015/025 | Batch 000/383 | Cost: 0.4743  
Epoch: 015/025 | Batch 120/383 | Cost: 0.4527  
Epoch: 015/025 | Batch 240/383 | Cost: 0.3980  
Epoch: 015/025 | Batch 360/383 | Cost: 0.4704



Epoch: 015/025 Train Acc.: 87.51% | Validation Acc.: 75.80%  
Time elapsed: 19.27 min  
Epoch: 016/025 | Batch 000/383 | Cost: 0.3573  
Epoch: 016/025 | Batch 120/383 | Cost: 0.5161  
Epoch: 016/025 | Batch 240/383 | Cost: 0.8183  
Epoch: 016/025 | Batch 360/383 | Cost: 0.5599  
Epoch: 016/025 Train Acc.: 84.63% | Validation Acc.: 73.50%  
Time elapsed: 20.55 min  
Epoch: 017/025 | Batch 000/383 | Cost: 0.3803  
Epoch: 017/025 | Batch 120/383 | Cost: 1.0083  
Epoch: 017/025 | Batch 240/383 | Cost: 0.6728  
Epoch: 017/025 | Batch 360/383 | Cost: 0.5373  
Epoch: 017/025 Train Acc.: 87.59% | Validation Acc.: 75.90%  
Time elapsed: 21.84 min  
Epoch: 018/025 | Batch 000/383 | Cost: 0.3428  
Epoch: 018/025 | Batch 120/383 | Cost: 0.3178  
Epoch: 018/025 | Batch 240/383 | Cost: 0.3678  
Epoch: 018/025 | Batch 360/383 | Cost: 0.4563  
Epoch: 018/025 Train Acc.: 91.10% | Validation Acc.: 75.30%  
Time elapsed: 23.13 min  
Epoch: 019/025 | Batch 000/383 | Cost: 0.2925  
Epoch: 019/025 | Batch 120/383 | Cost: 0.2200  
Epoch: 019/025 | Batch 240/383 | Cost: 0.3575  
Epoch: 019/025 | Batch 360/383 | Cost: 0.2210  
Epoch: 019/025 Train Acc.: 93.76% | Validation Acc.: 74.90%  
Time elapsed: 24.41 min  
Epoch: 020/025 | Batch 000/383 | Cost: 0.1384  
Epoch: 020/025 | Batch 120/383 | Cost: 0.3544  
Epoch: 020/025 | Batch 240/383 | Cost: 0.5099  
Epoch: 020/025 | Batch 360/383 | Cost: 0.2977  
Epoch: 020/025 Train Acc.: 93.00% | Validation Acc.: 77.70%  
Time elapsed: 25.69 min  
Epoch: 021/025 | Batch 000/383 | Cost: 0.2436  
Epoch: 021/025 | Batch 120/383 | Cost: 0.3992  
Epoch: 021/025 | Batch 240/383 | Cost: 0.6025  
Epoch: 021/025 | Batch 360/383 | Cost: 0.2882  
Epoch: 021/025 Train Acc.: 92.24% | Validation Acc.: 76.00%  
Time elapsed: 26.98 min  
Epoch: 022/025 | Batch 000/383 | Cost: 0.2165  
Epoch: 022/025 | Batch 120/383 | Cost: 0.2758  
Epoch: 022/025 | Batch 240/383 | Cost: 0.1705  
Epoch: 022/025 | Batch 360/383 | Cost: 0.3327  
Epoch: 022/025 Train Acc.: 95.40% | Validation Acc.: 77.40%  
Time elapsed: 28.26 min  
Epoch: 023/025 | Batch 000/383 | Cost: 0.1869  
Epoch: 023/025 | Batch 120/383 | Cost: 0.0938  
Epoch: 023/025 | Batch 240/383 | Cost: 0.2034  
Epoch: 023/025 | Batch 360/383 | Cost: 0.3387  
Epoch: 023/025 Train Acc.: 95.31% | Validation Acc.: 77.80%

```
Time elapsed: 29.54 min
Epoch: 024/025 | Batch 000/383 | Cost: 0.1274
Epoch: 024/025 | Batch 120/383 | Cost: 0.1440
Epoch: 024/025 | Batch 240/383 | Cost: 0.1691
Epoch: 024/025 | Batch 360/383 | Cost: 0.1815
Epoch: 024/025 Train Acc.: 96.22% | Validation Acc.: 76.30%
Time elapsed: 30.82 min
Epoch: 025/025 | Batch 000/383 | Cost: 0.0593
Epoch: 025/025 | Batch 120/383 | Cost: 0.0517
Epoch: 025/025 | Batch 240/383 | Cost: 0.1567
Epoch: 025/025 | Batch 360/383 | Cost: 0.1563
Epoch: 025/025 Train Acc.: 96.09% | Validation Acc.: 76.80%
Time elapsed: 32.12 min
Total Training Time: 32.12 min
```

```
with torch.set_grad_enabled(False): # save memory during inference
    print('Test accuracy: %.2f%%' % (compute_accuracy(model,
test_loader, device=DEVICE)))
```

```
Test accuracy: 74.04%
```